

Where Steaks Meet Statistics: Modeling Inventory Demand as a Finite-Horizon Markov Decision Process

Alex Huml¹, Carter Hall¹, Thomas Bridges¹, Ian Wilson¹

¹ Department of Statistics and Operations Research, University of North Carolina at Chapel Hill

1 Abstract

This report presents a practical intersection of statistics, economics, and industry by modeling inventory decisions for a local steakhouse, Stoney River Steakhouse and Grill. After an arduous process of converting physical data into a digital format, recent sales records for four main varieties of steak were utilized as respective foundations for finite-horizon Markov Decision Processes (MDPs). Prices were transformed according to the inflation rate, such that future decisions would account for economic inflation (FRED, 2022). The resulting optimal policies are reflective of potential ordering practices for the business, as they aspire to effectively meet customer demands, now and in the future.

2 Background

In aspiring to model these inventory decisions for Stoney River through an MDP, careful assumptions and observations had to first be made to ensure the relevance of any concluded results. At this location, the restaurant orders from tenderloin, bone-in ribeye, or “cowboy,” strip, or boneless ribeye loins/steaks. Additionally, Stoney River orders ground chuck meat as well, though as this is sometimes used in-place of other steaks as “side items,” the modeling process ignores these orders and considers the four aforementioned categories.

2.1 Important Observations

For each type of loin, of major importance is the amount Stoney River is capable of storing at any one time; there can be 84 tenderloins, 9 cowboy

loins, 12 strip loins, and 24 ribeye loins in Stoney River's storage. Such large figures allowed for the removal of any spoilage considerations, as demand for the steaks is fast enough in comparison such that meat seldom, if at all, spoils. As for the MDPs, these values, when converted to pounds, were able to serve as maxima for each process' respective state space.

Another complication arises from the explicit measurements of steaks as ordered by the customer; to circumnavigate this, all data was converted into pounds such that the end-result could be interpreted into pounds and partitioned into individual orders as the restaurant desires. For context, the following information is given for each loin:

- Strip: 5 loins per case, at an average of 13 lbs per loin
- Cowboy: 3 loins per case, at an average of 18 lbs per loin
- Ribeye: 4 loins per case, at an average of 15 lbs per loin
- Tenderloin: 12 loins per case, at an average of 7 lbs per loin

It should be noted that, as not all cows are created equal (differing in size by the season), formatting of the results in this manner allows for a better frame of reference when the restaurant places order(s) from their supplier(s).

2.2 Assumptions

Before presenting the framework behind the MDPs, a few important assumptions and considerations should be presented, namely:

1. The restaurant previously sold a 12-ounce filet, but has since removed the option from the menu, leaving an already existent option of ordering a 10-ounce instead. To account for the same demand, the assumption is made that customers who previously ordered the 12-ounce filet will now order its smaller counterpart.
2. Some steaks will inevitably not be to customers' satisfactions, and will therefore be thrown out. However, as the meat was utilized by the kitchen, the demand is recognizable from the restaurant's perspective, and so inflated demands due to mis-cooked steaks are assumed to be negligible.
3. When entering the data into Excel spreadsheets, an issue arose with respect to the classification of orders based on the week of the year. The assumption of every month having four weeks was made, wherein this last week could have any additional days that would otherwise

constitute a fifth. Hence, “4.1” in the Excel documents, for example, corresponds directly to the first week of April.

4. In actuality, the restaurant could order anywhere from one to four times per week. Given that this is unpredictable from the perspective of the authors (and a Poisson process in general that we have no data to quantify), the assumption is made that prices do not change mid-week, hence the data is recognized as week-long totals.
5. As no spoilage is assumed due to capacity restrictions, salvage costs are considered to be negligible.
6. To discretize the normal distributions associated with demand patterns for each type of steak, Equation 1 demonstrates the assumption for an arbitrary normal random variable, X . Furthermore, the normal distributions explained in Section 2.3 are assumed to be independent from one another, as there is no real and effective manner in which quantifying how an order of one steak could detract from the order totals of another.

$$F(x) = P(X = \lceil x \rceil) - P(X = \lfloor x \rfloor) \quad (1)$$

2.3 Creating Demand Probability Distributions

Unfortunately, Stoney River was not able to supply week-long data as desired, but instead gave year-long totals for the amount of each steak sold. This forced the creation of an accurate probability distribution that did not rely on the fitting of a model to data points. After consultation with the restaurant’s general manager, the conclusion was reached to use approximation to a normal distribution for each steak type. This decision was fortified due to assurance that the week-to-week demand was not a volatile figure, with the exception of holiday weeks or unusual “down periods.”

As the order data was in-hand, the standard deviation of the pounds of loins ordered per week was derived and used as a proxy for the standard deviation of any demand distribution. However, because orders are placed multiple times per week and not every pound ordered goes into making the steaks whose demands are being modeled, only half of the standard deviation was elected to be used to account for such a trend. Therefore, the following probability distributions were derived for each steak:

- Strip: $\mathcal{N}(112, 43)$
- Cowboy: $\mathcal{N}(54, 26)$
- Ribeye: $\mathcal{N}(119, 54)$
- Tenderloin: $\mathcal{N}(188, 89)$

2.4 Prices

As is the case with food, two prices are of importance: order price and selling price. The process to create the figure for the selling price was rather straightforward; as Stoney River sells different sized entrees of the same cut of beef (i.e. 10 and 16-ounce strips), computing a weighted average of the prices based on how much each steak was sold, in relation to the total, led to a single number representing price per pound for each steak.

The purchase prices for each steak type are presented in Figure 1, with the price per pound for each steak in the caption.

To model expected prices, price data was adjusted according to the rate of inflation. For each steak type, a scatter-plot relating price by week was matched to changes in the inflation rate. Expert predictions for inflation rates for the rest of the year were consulted, though of major relevance is the Russia-Ukraine conflict that is driving inflation rates up in the short-term. Equation 2 represents the price at an epoch, $Pr(t)$, modeled for each steak type in terms of average price over 2021, A . Note that A is given as follows for each steak type (all in USD) Strip: 8.96, Cowboy: 11.32, Ribeye: 12.68, Tenderloin: 14.98.

$$Pr(t) = A * (1 + \frac{0.05}{\sqrt{t} - 0.5}) \quad (2)$$

The evaluation behind this projection is that prices will decrease in the latter half of 2022, though remain higher than A .



Figure 1: Purchase prices for each steak, by week. The selling price, per pound, for each steak is as follows: *Strip* = 44 USD, *Cowboy* = 36 USD, *Ribeye* = 48 USD, *Tenderloin* = 75 USD.

3 Setup of the MDPs

As a summary of the presented background information, this section presents the information for the finite-horizon Markov Decision Processes, who in general reflect the 5-tuple

$$(T, S, A, p_t(s, a, j), r_t(s, a, j))$$

where $T = \{1, 2, \dots, 32\}$ represents the set of decision epochs corresponding to each week, since the beginning of May. Note the selling prices from Figure 1 are given as S_p . To avoid redundancy in this report, let s_{max} be the maximum state attainable for each steak type, given as:

- Strip: 168
- Cowboy: 171
- Ribeye: 360
- Tenderloin: 588

This gives the following for the state and action spaces S and A_s , respectively,

$$\begin{cases} S = \{0, 1, \dots, s_{max}\} \\ A_s = \{0, 1, \dots, s_{max} - s\} \forall s \in S \end{cases}$$

Additionally, the reward and transition probability functions are qualitatively equivalent, and so they are given universally by Equations 3 and 4. Note CDF represents the CDF of a normal distribution, using the normal distribution parameters presented in Section 2.3.

$$r_t(s, a, j) = \begin{cases} S_p * (s + a - j) * F(s + a - j) - a * Pr(t) & j \in [0, s + a] \\ S_p * (s + a) * (1 - \text{CDF}(s + a)) - a * Pr(t) & \text{else} \end{cases} \quad (3)$$

$$p_t(s, a, j) = \begin{cases} 0 & j > s + a \\ F(s + a - j) & 0 < j \leq s + a \\ 1 - \text{CDF}(s + a) & j = 0 \end{cases} \quad (4)$$

4 Results

4.1 Total Expected Rewards

Upon running the code presented in Appendix A (see the GitHub link included for a quicker implementation than sequential programming, as run-times were multiple hours in testing), Table 1 gives maximum total expected rewards, computed in epoch $t = 0$ as the MDP process begins with boundary conditions and works backwards, and the state in which these maximum rewards were attained. The purpose of including these figures is to indicate the best starting inventories.

	Strip Loin	Ribeye	Cowboy	Tenderloin
$\max_{s \in S} \{r_0(s)\}$	105,049.35	110,219.15	33366.79	279,138.14
$\arg \max_{s \in S} \{r_0(s)\}$	168	360	171	588

Table 1: Maximum Total Expected Rewards for each MDP process with the accompanying states (inventories) in which these were attained.

Not surprisingly, the tenderloin ascertains the highest maximum total expected reward across all steak types, with the ribeye, strip loin, and cowboy steaks ranked behind in decreasing order. A notable observation is that this ranking follows the profits per pound for each steak type, as well as this maximum reward unsurprisingly occurring when inventory is maximized. However, there is a caveat to the aforementioned rewards indicated by the optimal policies; in every case except that of the strip loin, the optimal policy calculated by the MDP never decides to order such that capacity is reached. Therefore, the “maximum-attainable total expected reward,” denoted by $\max_{s \in S^*} (r_0(s))$, where S^* is the subset of attainable states, is derived by using the maximum expected total reward for the highest state (i.e. inventory) attainable for a particular MDP. The findings in this context are presented in Table 2.

	Strip Loin	Ribeye	Cowboy	Tenderloin
$\max_{s \in S^*} \{r_0(s)\}$	105,049.35	109,109.90	32,927.01	276,843.92
$\arg \max_{s \in S^*} \{r_0(s)\}$	168	277	134	442

Table 2: Maximum-Attainable Total Expected Rewards for each MDP process with the accompanying states (inventories) in which these were attained.

It is worth noting that while the tenderloin has the highest maximum reward (and maximum-attainable reward), it has the highest return for purchase price. At a *sale-to-purchase ratio* of 5.007, it slightly exceeds that of the strip loin at 4.91 (though this may indicate why the inventory for the strip is consistently filled – see Section 4.2 for more detail). As the tenderloin essentially becomes Stoney River’s “cash cow” (as it *is* a steak), the restaurant unsurprisingly sets aside a large portion of its storage for purchasing tenderloins.

In fact, the sale-to-purchase ratio is minimized by the cowboy steak, which unsurprisingly has the lowest capacity for storage within the restaurant.

4.2 Optimal Policies

Interestingly, optimal policies did not follow conventional logic, namely to simply fill inventory completely in every decision epoch. Whether this was a direct consequence of inflation or potentially waning demand for particular steak types, the results dictated a varying decision process for each steak, with exception of the strip loin. Perusal of the optimal policies is fully at the reader’s discretion, though decision rules for Week 1 (the first decision epoch, $t = 1$) will be presented here as an indicator of how to effectively read the policies. Additionally, the repository linked in Appendix A will have functionality where readers may choose to more programmatically look at the optimal policies for any given week.

Given the general framework for a decision rule, d_t , as

$$\begin{cases} d_t = \Sigma_t - s & s < \sigma_t \\ 0 & \text{otherwise} \end{cases} \quad (5)$$

for the first epoch, at time $t = 1$, the following values for Σ_t, σ_t are given for each steak type:

- Strip Loin : $\Sigma_t = \sigma_t = 168$
- Cowboy : $\Sigma_t = \sigma_t = 134$
- Ribeye: $\Sigma_t = \sigma_t = 278$
- Tenderloin: $\Sigma_t = \sigma_t = 442$

For all other decision epochs (with the exception of Week 32, as the setup dictated no ordering to occur on that week and any salvage costs, which were assumed to be negligible, to be evaluated), the ordering policy varies by week and by steak. The general trend discovered by the ordering policies

was that, in the case of the strip loin, Stoney River should continue ordering to-capacity while adhering to more unique policies in all other cases. For the future, however, this could suggest an increase in storage for strip loins, as if the steak is constantly being ordered, having increased amounts on-hand could directly translate to increased profits.

5 Conclusion

By implementing the presented policies as advised by Markov Decision Processes, Stoney River Steakhouse will be able to maximize their total expected rewards (i.e. increase profit) while meeting customers' demands and accounting for future inflation. The processes incorporated year-long demand data fitted to normal distributions and converted into usable, week-long data. Additionally, inflation rates were incorporated in adjusting the purchase prices of each steak. Such preprocessing was then implemented in a Python 3.9 class that housed methods to solve a finite-horizon MDP to optimality and return optimal ordering policies.

6 Literature Cited

"Beef2022 Data - 2001-2021 Historical - 2023 Forecast - Price - Quote - Chart." *Beef - 2022 Data - 2001-2021 Historical - 2023 Forecast - Price - Quote - Chart*, <https://tradingeconomics.com/commodity/beef>.

"Prices - Inflation Forecast - OECD Data." TheOECD, <https://data.oecd.org/price/inflation-forecast.htm>.

"10-Year Breakeven Inflation Rate." FRED, 29 Apr. 2022, <https://fred.stlouisfed.org/series/T10YIE>.

"Consumer Price Index for All Urban Consumers: All Items in U.S. City Average." FRED, 12 Apr. 2022, <https://fred.stlouisfed.org/series/CPIAUCSL>.

7 Appendix A: Python Code and Excel Files

This appendix contains the relevant Python code and Excel outputs, beginning on the following page. The code, paper, and related output files can be found at the following link: <https://github.com/halljc76/STOR515FinalProj>.

```

class WeLoveSteak:
    """Modeling the demand for delicious steaks as a finite-
        horizon Markov Decision
        Process."""
    def __init__(self, M, N, mu, sigma, sellPrice, priceyint,
                  steakName):
        """Initialization Function.

        Parameters:
        M := Limit on storage capacity.
        N := Number of decision epochs.
        S := State Space. T := Epoch set.
        mu, sigma := Normal dist parameters.
        priceyint := P_0
        steakName := Name to pass into Excel writer.
        """
        self.M = M
        self.N = N
        self.S = set(range(self.M+1))
        self.T = set(range(1,self.N+1))
        self.mu = mu
        self.sigma = sigma
        self.sp = sellPrice
        self.priceyint = priceyint
        self.ustar, self.astar = self.optimality(steakName)
        self.storeResults(steakName)

    def A(self, s):
        """Function representing action space.

        Parameters:
        s := State from state space S.
        """
        return set(range(0,self.M - s + 1))
    def f(self, s):
        """Revenue generated from selling loins at sp/pound.

        Parameters:
        s := State indicating number of loins sold.
        """
        return self.sp * s
    def price(self, t):
        """Price for Strip Loin, adjusted for inflation via
            CPI metric.

        Parameters:
        t := Epoch in T
        """

```

```

        return self.priceyint * (1 + (0.05 / (np.sqrt(t) - 0.5)))

def CDF(self, s):
    """Function for discretized probability that uses the
        normal CDF.

    Parameters:
        s := State in S
    """
    return stats.NormalDist(self.mu, self.sigma).cdf(s)

def pt(self, j, s, a):
    """Transition probability function.

    Parameters:
        j,s := states in S
        a := action in A

    Returns:
        0 if j > s + a
        CDF(s + a - j) - CDF(s + a - j - 1) if 0 < j <= s + a
        1 - CDF(s + a) otherwise
    """
    assert type(j) == type(s) == type(a) == type(3)
    return 0 if j > s + a else
    ((self.CDF(s + a - j) - self.CDF(s + a - j - 1))
     if ((0 < j) and (j <= s + a))
     else (1 - self.CDF(s + a)))

def rt(self, t, s, a = 0):
    """Reward function for finite-horizon MDP setup.

    Parameters:
        t := Epoch in T
        s := State in S
        a := Action in A
    """
    if t == self.N:
        return 0
    elif s+a == 0:
        return 0
    elif s+a == 1:
        return (1-self.CDF(1))*self.f(1)
    er = 0
    er += (1-self.CDF(s+a))*self.f(s+a)
    for i in range(s+a-1):
        er += (self.f(i) * (self.CDF(i)-self.CDF(i-1)))
    return er - self.price(t)*a

```

```

def optimalAction(self, d):
    '''Function to determine which action corresponds to
        optimal policy.

        Parameters
        -----
        d: dict -> dictionary of (action, optimality eqn.
            value)
    '''
    v = max(d.values())
    for k in d.keys():
        if d[k] == v:
            return k

    raise ValueError("Should never get here, but.")
def optimality(self, steakName):
    '''Implements algorithm to find optimal policy via
        backward induction.'''
    t = len(self.T) # From 1 to N + 1, this returns N
    ustar = np.zeros([len(self.S), t])
    astar = np.zeros([len(self.S), t])

    # BC computation --  $u_{N+1}^{*}(s) = r_N(s)$  for all  $s$  in  $S$ 
    for s in self.S:
        ustar[s,-1] = self.rt(t,s)

    while t != 1:
        print(steakName + ": " + str(t))
        t -= 1
        for s in self.S:
            l = dict()
            for a in self.A(s):
                temp = self.rt(t,s,a) + sum([self.pt(j,s,
                    a) * ustar[j,t] for
                    j in self.S])

                l[a] = l.get(a, 0) + temp
            ustar[s,t-1] = round(max(l.values()), 4) #
                                                    Rounding this
                                                    because of PDF
                                                    output --
                                                    doesn't change
                                                    answer
            astar[s,t-1] = self.optimalAction(l)

    ustar = self.getOptimalityTable(ustar)
    astar = self.getOptimalPolicy(astar)

```

```

    return ustar, astar
def getOptimalPolicy(self, astar):
    '''Pretty-prints the optimal policy dataframe.'''
    return pd.DataFrame(astar,
        columns = ['Week {}'.format(t) for t in
                    range(1,self.N+1)],
        index    = ['State {}'.format(s)
                    for s in range(0,len(self.S))])

def getOptimalityTable(self, ustar):
    '''Pretty-prints the total expected reward dataframe.
        ,,,
    return pd.DataFrame(ustar, columns =
        ['Week {}'.format(t) for t in range(1,self.N+1)],
        index    = ['State {}'.format(s)
                    for s in range(0,len(self.S))])

def storeResults(self, steakName):
    with pd.ExcelWriter("./MDPResults" + steakName + ".
                        .xlsx") as writer:
        self.astar.to_excel(writer, sheet_name="ExpReward
                                ", index=True)
        self.ustar.to_excel(writer, sheet_name="OptPolicy
                                ", index=True)

    return

```